

NIXORA AI

Construction Intelligence Platform

A complete engineering proposal, phased delivery plan, and milestone-based estimate (INR) to take NIXORA from concept to a **100% production-ready, revenue-generating B2B SaaS platform** — not a pilot, not an MVP that stalls. Built on the same stack and architecture that powers our own SaaS products.

Full platform **₹1.30 Cr – ₹1.77 Cr**

Commercial MVP in **~12 weeks**

Full GA in **4–6 months**

Stack **Go · Temporal · React · PostgreSQL**

CLIENT	ENGAGEMENT	RECURRING	DELIVERY PARTNER
NIXORA AI	Milestone fixed-bid	Support + hosting AMC	DevzStudio
STACK BASIS	Proven in production		

From budget overruns to a Construction Operating System

NIXORA AI helps construction companies **prevent budget overruns, reduce project risk, and improve profitability** by turning the documents they already have — invoices, POs, PDFs, spreadsheets — into live financial and operational intelligence, with an AI copilot answering plain-language questions like *"which projects are over budget?"*

We propose building this in **five evolving releases**, each independently valuable and shippable, on a shared multi-tenant platform with billing, per-seat licensing, AI-usage metering, and audit engineered in from day one — the commercial machinery that turns a great product into a great *business*.

₹1.52 Cr

Recommended planning budget for the full platform (range ₹1.30–1.77 Cr, ex-GST).

9

Payment-linked milestones, each with a hard acceptance gate & demo.

~12 wk

To a billable commercial MVP (Phase 1) on the standard track.

4–6 mo

To full GA — all 5 phases, hardened & production-ready (standard / fast-track).

Why this is low-risk: the hardest parts of NIXORA — multi-tenant auth, role-based access, usage metering, billing, audit logging, document ingestion, webhooks & automation — are subsystems we already run in production across our own SaaS products. For NIXORA they are **proven patterns on a proven stack, not R&D**. That means predictable delivery, fewer surprises, and engineering decisions that have already survived real-world load.

We don't theorise about B2B SaaS — **we operate it**

DevzStudio — *"Technology Right from Scratch"* — is a senior, product-minded engineering studio. We ship and run our own SaaS products, so we estimate from real operational experience, not slideware.



IntlPull

intlpull.com

AI-first Translation Management System with native MCP integration — translate from Claude, Cursor & VS Code. Multi-tenant orgs, billing, usage metering, RBAC & audit in production.

100+ Startups → Public sector

organizations incl. Nava PBC — US government healthcare programs (in trial)



UseLoyalty

useloyalty.app

AI-powered loyalty engine for rewards, missions, tiers, referrals and retention campaigns — event-driven rules & analytics at scale.



Ekada

ekada.app

Whitelabel commerce SaaS — storefronts, reservations, invoices & QR menus, with merchant-connected payment gateways. Multi-tenant, white-labelled, billing-aware.



ClixWall

clixwall.com

Website monetization & advertising marketplace.



LootGO

lootgo.app

Real-world treasure hunt turning daily steps into rewards — **87.6K+ followers on X**, a fast-growing global userbase & brand partnerships.



beyondClub

beyondclub.xyz

No-code web3 loyalty & membership for brands & creators.



CheckIn

checkincaster.xyz

Geolocation 'check-in & earn' social primitive on Farcaster.



HippoDash

hippodash.com

Analytics dashboard that turns scattered data into clear, actionable insights.



StoryForU

storyforu.com

AI-powered personalized children's stories with generated illustrations.



BeaverGrow

beavergrow.com


Customizable productivity platform — tasks, notes & health tracking.



CodeKeep

codekeep.io

Save & organise code snippets for developers — **#1 Product of the Day** on Product Hunt.



Product thinking

We design for the business model — pricing, limits, margins — not just features.



Secure by default

Input validation, parameterised queries, HttpOnly sessions, OWASP-aligned reviews.



Real AI/LLM ops

RAG, document AI, Google Cloud Vision, and cost-controlled multi-model LLM gateways in production.



Scale-ready

No N+1 queries, eager loading, observability, and load-tested release gates.

Partnering with the best in tech

We build alongside — and are trusted by — the platforms and firms our clients rely on.

Palo-IT

OpenAI

Anthropic · Claude

GitHub

Microsoft

AWS

Google Cloud

MongoDB

Visual Studio


Bain & Company

Freedom of Mobility Forum

Engineering partner to **Palo-IT** — a global Agile software consultancy — across **API services** & **web & app development** engagements.

Built to scale — proven on one app

A single consumer app we built scaled to **74M+ ad impressions** and **\$16.5K+** in ad revenue — a real measure of the traffic & growth our architecture is engineered to handle.



74M+ **\$16.5K+**

impressions *ad revenue*

162M+

ad requests

One app · AdMob lifetime —
\$16,482.69 earnings · 74,011,401
impressions · 162,398,961
requests · 566,153 clicks.

Five releases. One platform. Compounding value.

Each phase is a real, shippable release with its own ROI — so NIXORA can demo, sell, and gather feedback continuously rather than waiting for a "big bang."

PHASE	THEME	HEADLINE CAPABILITIES	OUTCOME FOR THE CUSTOMER
PHASE 1	Budget & Risk MVP	Doc upload (PDF/Excel/invoice/PO) → AI extraction & categorisation, planned-vs-actual, variance, budget-overrun alerts, anomaly & risk scoring, exec dashboard, AI copilot, PDF reports	See project financial health & detect budget issues with no manual data entry
PHASE 2	Project Intelligence	Timelines & milestones, schedule AI (delay prediction, critical-path warnings, completion forecast), project health score, workforce/resource planning, notifications	Predict delays & resource conflicts before they hit the schedule
PHASE 3	Material Intelligence	Material & supplier DB, ordered/delivered/used tracking, waste & overconsumption detection, reorder & price-trend AI, supplier comparison & alternatives	Control material spend, cut waste, optimise procurement
PHASE 4	Construction Copilot	Advanced AI assistant ("why is Project A losing money?"), predictive analytics (budget/profit/cash-flow/completion), prescriptive cost-saving & procurement recommendations	An AI analyst that explains <i>why</i> and recommends <i>what to do</i>
PHASE 5	Construction OS	Multi-project portfolio intelligence, site-photo upload + AI image/progress analysis, mobile field app, integrations (accounting/ERP/procurement/scheduling)	The central intelligence hub for the whole company

*Build-first → Phases 1 + platform core. Build-later → Phases 2–4. Build-last → Phase 5 (full OS).
This sequencing matches NIXORA's own roadmap priorities.*

Built on the same stack that powers our SaaS

NIXORA is engineered on the identical, production-hardened stack and platform architecture we run at IntlPull and our other products. Every core subsystem below is something we have built, shipped, and operate at scale today — so for NIXORA they are dependable engineering, not experiments.

SUBSYSTEM	ENGINEERED FOR NIXORA	TRACK RECORD
Multi-tenant orgs + RBAC (owner/admin/member/viewer)	Tenancy isolation, role & permission model	In production
Auth — API keys + JWT/JWKS sessions, SSO	Secure auth for web, mobile & server-to-server	In production
Billing — subscriptions, checkout, invoicing	Plans, proration, invoicing, dunning	In production
Usage metering & limit enforcement	Real-time meters, soft/hard caps per org	In production
Audit / event log	Immutable activity trail, exportable	In production
Webhooks + automation rules engine	Event-driven alerts & integrations	In production
Document upload + parse / import pipeline	Ingestion for invoices, POs, Excel, PDFs	In production
AI/LLM & workflow layer	Temporal-orchestrated extraction, RAG copilot, forecasting, Vision	In production

Each subsystem is fully engineered, tested and delivered as part of NIXORA — but because the patterns are battle-tested in our live products, you get a **de-risked build with predictable timelines and production-grade quality from day one.**

How it's engineered

A modern, cloud-native, multi-tenant architecture with a dedicated AI services layer and a metered-billing core — the same architecture that powers our production SaaS. Strong typing, no any, no N+1 queries, secure by default. The core API is built in **Go (Golang)** for fast, low-latency performance, with **Temporal** for durable, crash-safe handling of long-running workflows.

① **Client layer** – Manager web dashboard + mobile-friendly field view + native mobile app (Phase 5)

Next.js / React

TypeScript

Tailwind + shadcn/ui

React Native / Expo

② **API & platform core** – high-performance REST API in Go (Chi), RBAC, tenancy, audit, webhooks, automation rules

Go (Golang)

Chi router

Auth: JWT + API keys

Rate limiting

③ **AI services layer** – document extraction (OCR + LLM), RAG copilot, risk scoring, forecasting, site-photo vision – orchestrated with Temporal durable workflows

Temporal workflows

Google Cloud Vision / Document AI

LLM Gateway (Gemini · Claude · GPT)

RAG / embeddings

pgvector

Usage metering

④ **Monetization core** – subscriptions, per-seat, AI-usage ledger, metering, limit enforcement, invoicing

Stripe

Pluggable gateways

Invoicing

Dunning

⑤ **Data & infra** – relational data, object storage, queues, observability, backups

PostgreSQL

Redis

S3 / object storage

Temporal / workers

Sentry + metrics

AWS / GCP

⑥ **Integrations (Phase 5)** – accounting, ERP, procurement, scheduling

Webhooks

OAuth connectors

Import/Export

We build the engine — **you choose the model**

We engineer the full commercial machinery — subscriptions, per-seat licensing, AI-usage metering, limit enforcement, invoicing and audit. **How you price and package it is your commercial decision**; the platform supports whichever model NIXORA chooses, and these are the key open questions to settle together.



Per-seat licensing

Base plan + priced seats, role-aware. The engine enforces seat counts; tier & price points are yours to set.



AI-usage credits

Meter each AI action (parse, copilot answer, forecast, vision) as billable units. You decide inclusions, overage & pack pricing.



Limit enforcement

Real-time usage meters per org with soft caps (notify) and hard caps (block / auto-purchase) — configurable.

Open questions for NIXORA

- Per-seat vs flat vs hybrid pricing — and your target ARPU?
- How many AI credits bundled per plan, and overage behaviour?
- Top-up packs & volume discounts — sold, or usage-billed monthly?
- Free trial / freemium tier, or sales-led only?
- Annual vs monthly billing & discount strategy?

Reference usage weighting

How different AI actions can map to credits — an illustration to frame the conversation. Final packaging & pricing are NIXORA's to set.

AI ACTION	CONSUMES
Parse 1 document page	1.0 credit
Copilot answer (RAG)	0.5 credit
Forecast / analysis run	2.0 credits
Site-photo vision analysis	1.5 credits

Relative weighting only — heavier AI actions consume more credits. DevzStudio handles the underlying AI provisioning.

Built in **Milestone M2** — including payment-gateway integration (Stripe and other supported gateways), invoicing and automated dunning for failed payments.

Nine payment-linked milestones

Every milestone ends in a working demo and a written acceptance gate. Payment is released only on acceptance — so risk stays balanced for both sides.

#	MILESTONE	KEY DELIVERABLES	COST (₹ LAKH)	PAY %
M0	Discovery, Architecture & Design System	Requirements, data model, security model, UI kit, clickable prototype	8 – 11	10%
M1	Platform Core	Multi-tenant, Auth/SSO, RBAC, company/team, immutable audit log	15 – 20	12%
M2	Monetization & Billing Engine	Subscriptions, per-seat, AI-usage ledger, metering, limits, Stripe + supported gateways, invoicing, dunning	11 – 15	10%
M3	Commercial MVP — Phase 1	Doc AI extraction, budget intelligence, risk detection, exec dashboard, AI copilot, PDF reports	28 – 37	24%
M4	Project Intelligence — Phase 2	Timelines, schedule AI, health score, workforce, notifications	15 – 20	12%
M5	Material & Supplier Intelligence — Phase 3	Material DB, tracking, waste AI, supplier comparison	12 – 16	9%
M6	Advanced Copilot & Predictive — Phase 4	Advanced assistant, forecasting, prescriptive recommendations	13 – 18	9%
Total — full platform (ex-GST)			130 – 177	100%

#	MILESTONE	KEY DELIVERABLES	COST (₹ LAKH)	PAY %
M7	Construction OS — Phase 5	Portfolio intelligence, site-photo vision AI, mobile app, ERP/accounting integrations	20 – 28	10%
M8	Hardening, Security Audit & GA Launch	Load/perf testing, security remediation, observability, docs, 100% production GA	8 – 12	4%
Total — full platform (ex-GST)			130 – 177	100%

*M0+M1+M2+M3 = **commercial MVP gate** (billable, multi-tenant, Phase 1 AI) ≈ ₹62–83 L. M8 = **100% production GA**, all phases hardened.*

Three commercial gates

GATE A

Commercial MVP

₹62 – 83 L

M0–M3. A billable, multi-tenant Phase-1 product you can sell and onboard real customers onto.

GATE B

Intelligence Platform

₹90 L – 1.18 Cr

Through M5 — adds Project & Material intelligence. The product most customers will pay a premium for.

GATE C

Full Construction OS

₹1.30 – 1.77 Cr

All 9 milestones, hardened, mobile + integrations, 100% production GA. Recommended planning figure **₹1.52 Cr**.

All figures ex-GST. GST @ 18% applies on services. Fast-track delivery adds approximately **15–20%** to milestone costs (additional engineers for maximum parallelism — see Timeline). Estimates assume the scope described here; material scope changes are handled via the change process in Section 11.

How fast — and when it's 100% done

MVP	FULL GA	TEAM
~12 wk	~6 mo	~9–11

Standard track · 2 parallel pods

M0 Discovery · 2.5w



M1 Platform + M2 Billing (parallel) · 4w



M3 MVP (Phase 1) · 6w → MVP ~wk 12



M4 + M5 intelligence (parallel) · 5w



M6 + M7 Copilot + OS (parallel) · 6w



M8 Harden + GA · 2.5w → GA ~month 6



MVP	FULL GA	TEAM
~8 wk	~4 mo	~12–14

Fast-track · 3 parallel pods (+15–20% cost)

M0 Discovery · 1.5w



M1 + M2 (parallel) · 3w



M3 MVP (Phase 1) · 4w → MVP ~wk 8



M4 + M5 (parallel) · 3.5w



M6 + M7 (parallel) · 4.5w



M8 Harden + GA · 2w → GA ~month 4



Fast-track runs more phases concurrently with a larger team. It compresses calendar time ~30% further; the premium covers the extra engineers and coordination.

"When is it 100% working — not a pilot?" Production GA is reached at the end of **M8**: ~**6 months** (standard) or ~**4 months** (fast-track) from kickoff. **Pace scales with team size and budget** — a larger parallel pod compresses delivery, so we right-size the team to NIXORA's target launch date and budget. At GA every phase is feature-complete, security-audited, load-tested, documented, and live for paying customers.

Who builds it

A senior, cross-functional pod led by a principal engineer — scaled up for fast-track delivery.

Principal Engineer

Architecture, AI design & code review

Backend · Go / Temporal

APIs, billing core & durable workflows

Frontend · React / Next

Dashboard, copilot UI & reports

AI / ML Engineer

Extraction, RAG, forecasting & vision

Mobile · React Native

Field app for on-site teams

DevOps · QA · Design · PM

Infra, test automation, UX & delivery

One combined monthly plan — fully managed

After launch, DevzStudio keeps NIXORA running as a single managed engagement that **combines cloud hosting/infrastructure, monitoring, security patching and ongoing support** into one predictable monthly fee — so you have one partner accountable for uptime and one invoice to track.

STANDARD

Managed Care

₹1.1 L / mo

Hosting + infra + monitoring, bug fixes, security patches, dependency upgrades. Business-hours support · 99.5% uptime.

PRIORITY

Managed Care+

₹1.8 L / mo

Everything in Standard + extended-hours support, faster SLA, minor enhancements pool & proactive performance tuning · 99.9% uptime.

ENTERPRISE

Managed Care 24x7

₹2.8 L+ / mo

Everything in Priority + 24x7 on-call, dedicated engineer allocation, custom SLA & scaling support. Custom-scoped.

Plans scale with usage — heavy AI/LLM consumption, large data volumes or high traffic move the engagement up a tier. Tier is confirmed against NIXORA's expected scale before launch.

Out-of-scope work (new features beyond the agreed scope, third-party requests, major redesigns) is quoted and billed separately, priced according to complexity and effort, and only started after written approval.

Clear, standard & predictable

Payment schedule

- **25% advance** on signing (mobilisation) — adjusted across M0–M2
- Balance released **per milestone** on written acceptance
- Net-7 invoicing; milestone demo precedes each invoice
- **GST @ 18%** applicable extra on all service invoices
- TDS (if applicable) deducted as per Indian IT rules; certificates shared

Included & out of scope

- ✓ Source code + full IP transfer on final payment
- ✓ Documentation, deployment runbooks & handover
- ✓ 30-day post-GA warranty (defects fixed free)
- **Excl.** content / data entry & copywriting
- **Excl.** scope added after sign-off — quoted by complexity

Key assumptions & risks

Sample documents (invoices/POs/Excel) provided in M0 to tune extraction accuracy.

Timely feedback & acceptance within 3 business days per gate keeps the timeline on track.

LLM provider pricing may shift; the metering model surfaces and absorbs this transparently.

→ NEXT STEPS

Start with a paid Discovery (M0)

A 2–3 week, ₹8–11 L Discovery sprint de-risks the entire program: locked architecture, data model, security model, a clickable prototype, and a confirmed fixed quote for M1 onward — fully credited toward the build if you proceed.

A couple of things we'd like to confirm with NIXORA

- Do you have an existing **design system, brand kit, or a competitor / benchmark product** we should reference when shaping the features and UX?
- Can you give us access to a **construction subject-matter expert** — someone who can bridge the technical build with real-world site, project & estimation problems on the ground?



DevzStudio

Technology Right from Scratch · We build & operate B2B SaaS — intlpull.com · useloyalty.app · ekada.app · devzstudio.com